

Modelling Apple's M1 Firestorm Core

Processor Simulation with SimEng


Li-yun Wang

Supervisor: Simon McIntosh-Smith

2022 Sep

Apple's M1 Firestorm

- Apple Inc. introduced its self-designed M1 chip for MacBook in 2020.
 - In M1 CPU, 4 **high-performance cores, Firestorm**, and 4 high-efficiency cores, Icestorm.
- Firestorm gets impressive performance as compared with its competitor.
 - However, most of the details remain in a black box, disclosed by Apple.



<https://www.anandtech.com/show/16226/apple-silicon-m1-a14-deep-dive/2>, (2021).

<https://www.apple.com/uk/newsroom/2020/11/apple-unleashes-m1/>, (2020).

Outline

- **Background**
- **Aims & objectives**
- **Design & Implementation**
 - Configuration
 - Verification
- **Analysis & Evaluation**
 - Execution Units with Different Opcodes
 - ROB size & register file size
- **Future Work**
- **Conclusion**
- **Reference**

Background

□ Modern Processor microarchitecture

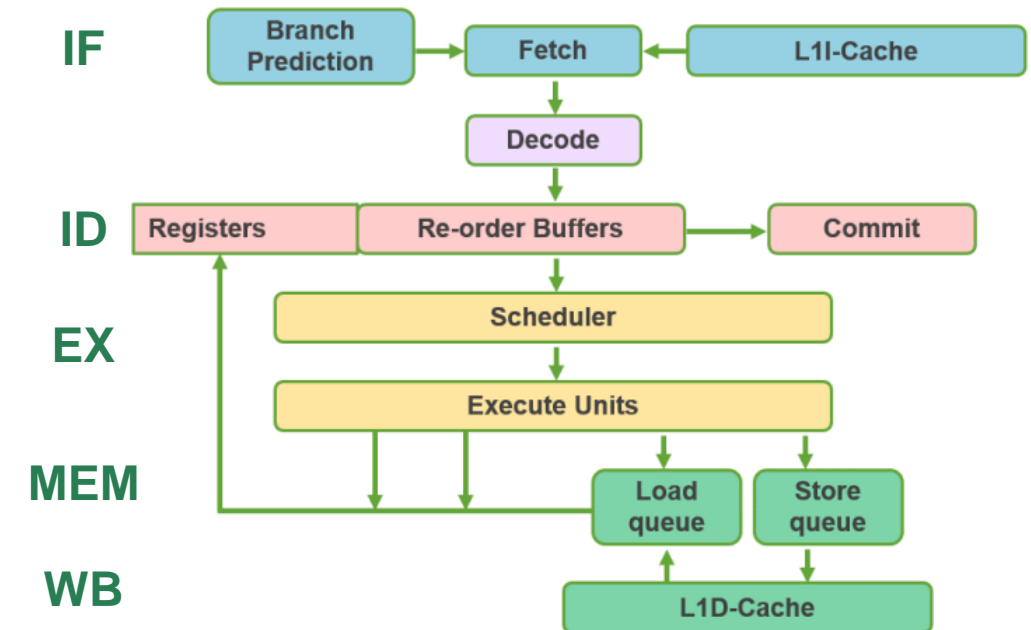
- 5-stage pipeline: IF -> ID -> EX -> MEM -> WB
- Out-of-order
- Superscalar

□ Processor simulation

- Crucial approach for processor design

□ SimEng simulation

- A framework developed by Simon's team
- Supporting sophisticated, superscalar, out-of-order microarchitectures
- Simple and efficient



<https://uob-hpc.github.io/SimEng/>, (2021).

D. A. Patterson and J. L. Hennessy, *Computer organization and design ARM edition: the hardware software interface*.

Aim & Objectives

□ Why we need a SimEng simulator for M1 Firestorm?

- M1 Firestorm: renowned, but only a few open source for relevant research
- SimEng simulation: simple, efficient and open source.

□ In this project,

1. Modelling Apple's M1 Firestorm core with SimEng simulation
2. Verifying the SimEng simulator
 - Demonstrate the features of Firestorm
 - Accuracy as the simulated cycles being within 10% of Firestorm
 - Benchmark with other's research

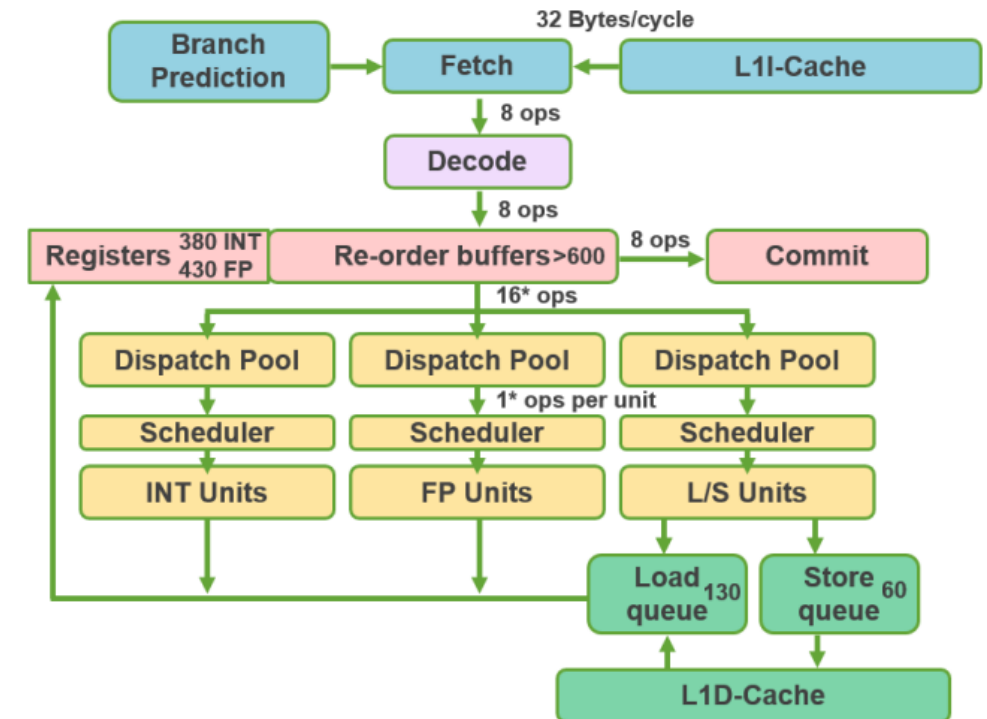
Design and Implementation

1. Configuration of SimEng simulation

- Core
- Bandwidth
- Buffer size:
 - Re-order buffer, register files & L/S queue
- Execution units:
 - latency, throughput & number of units

2. Verification

- Experiments & program with Assembly codes
 - Loops 1 million times to eliminate noise
- Perf tool: output the clock cycle of real machine




D. Johnson, <https://dougallj.github.io/applecpu/firestorm.html>, (2021).

Cardyak, https://drive.google.com/drive/folders/1-Ls0SiqXBCBhMGr87IdA_47GETWfjCzQ, (2022).

Execution Units - Latency

□ Latency

- **Processor's scalar performance**
 - Interval to the calculated data ready to other operations
 - Shorter, higher performance
- **Different latency of opcodes**
- **With configuration, SimEng can demonstrate these features.**
 - Accuracy: miss-matching <1%, goal <10%



K. Hwang and N. Jotwani, *Advanced computer architecture: parallelism, scalability, programmability*, 1993.


Execution Units – Counts & Throughput

□ Unit Counts

- Max. number of units can perform the operations

□ Throughput

- Processor's parallel performance
 - Total instructions completed in a certain time
 - Depending on the number of units
 - higher, higher performance
- Most cases SimEng performs perfectly; however, some finds larger miss-matching ratio >20%



* For add, 4 units with additional 2 units used by loop counter

Execution Units – Accuracy Summary

Accuracy of simulation

- Miss-matching ratio <3%, goal <10%
- Some indep. cases >30% → loop


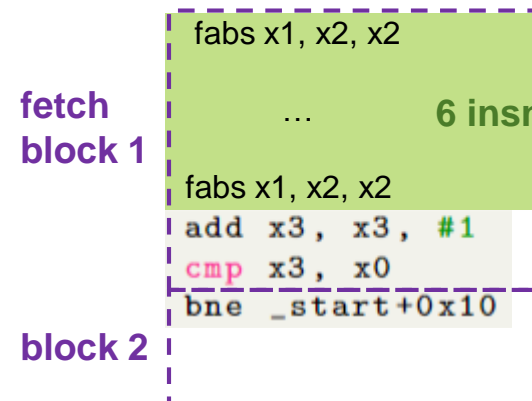
Loop prediction

- Firestorm w/ loop prediction
 - After hitting branch, fetch the loop body.
- SimEng ver. w/o loop prediction
 - As the loop body contains > 8 insns, SimEng takes at least 2 cycles.
 - Fetch rate dominates, instead of execution rate.

M. Kobayashi, *Dynamic characteristics of loops*, IEEE Transactions on Computers, 33 (1984), pp. 125–132.

Instruction Group List				
Instruction Groups	opcode	(A)dep.	(B)indep.	(C)indep.
INT_SIMPLE_ARTH_NOSHIFT	add	0%	<1%	0-48%
INT_SIMPLE_LOGICAL_NOSHIFT	and	0%	<1%	0-48%
INT_MUL	mul	0%	<1%	<1%
INT_DIV_OR_SQRT	sdiv	0%	<1%	<1%
FP_SIMPLE_ARTH_NOSHIFT	fabs	0%	<1%	0-33%
FP_SIMPLE_LOGICAL_NOSHIFT	bic	0%	<1%	0-33%
FP_MUL	fmul	0%	<1%	0-33%
FP_DIV_OR_SQRT	fddiv	0%	<1%	<1%
LOAD	ldr	<2%	<1%	0-80%
STORE	str	-	<1%	0-20%

Loop body:



ROB Size & Register File Size

- Buffers/queues in microarchitecture
 - Capable of performing superscalar execution
 - Larger, higher potential performance


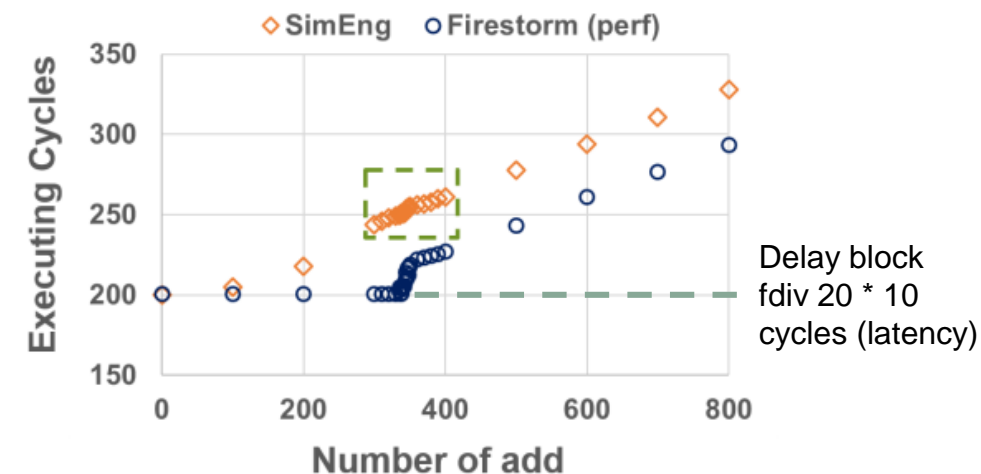
- Sizes of buffers

- Cycles jump as the buffer is filled

INT register file	FP register file	ROB
344	400	600

- SimEng can match the position of glitch
- However, SimEng gets almost linear curve impacted by fetch rate for loops

For INT register file size,

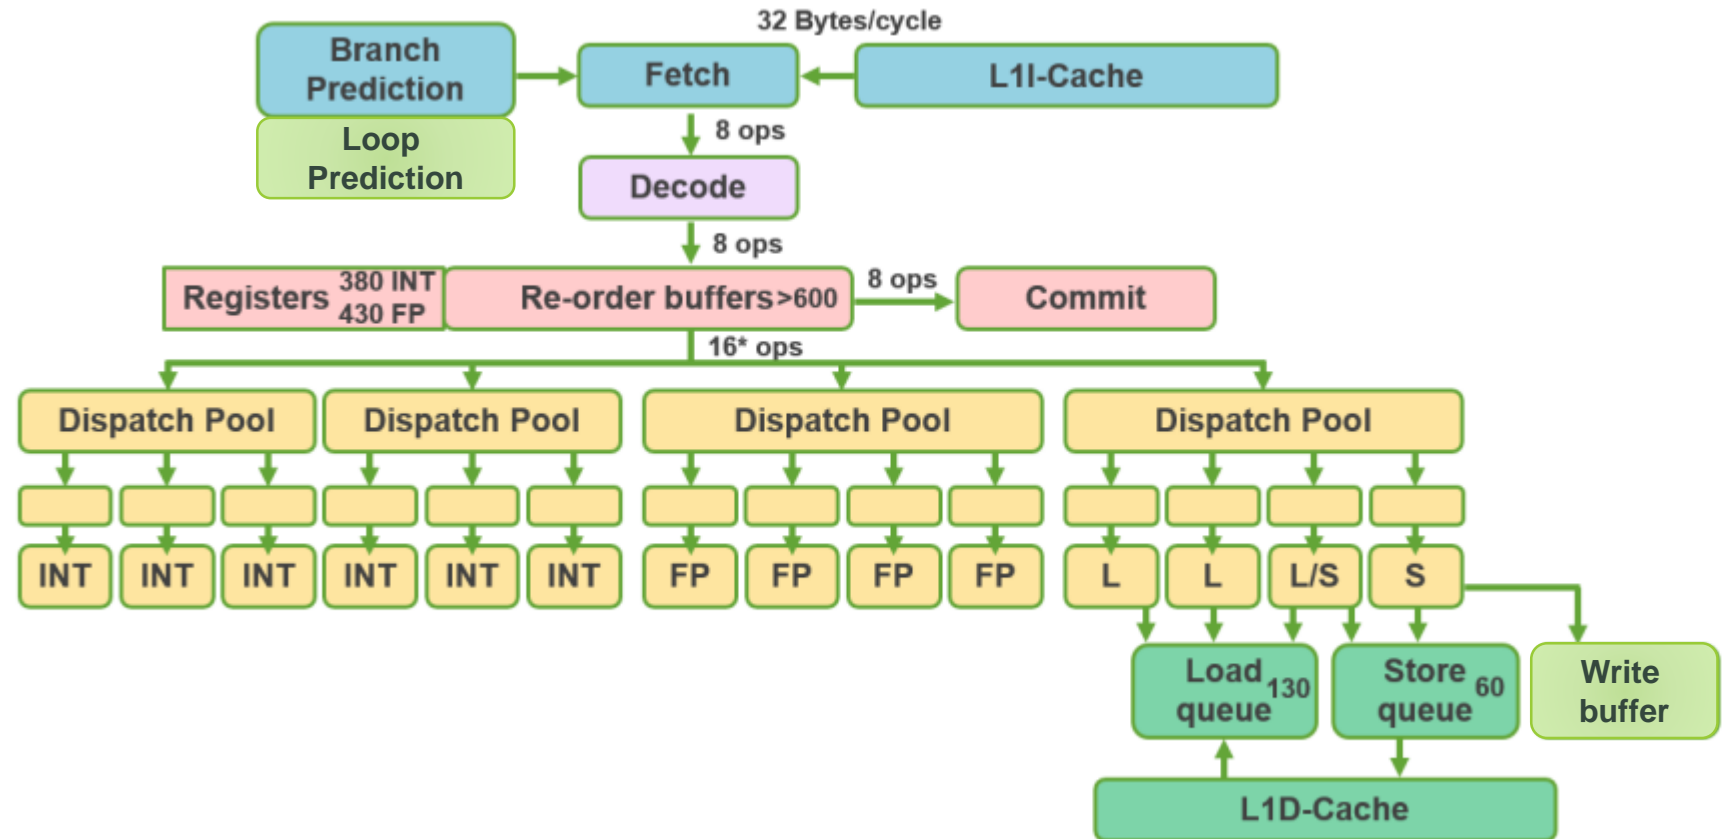


Future Work

□ Loop prediction

□ Execution units

□ Write buffer



M. De Alba and D. Kaeli, *Path-based hardware loop prediction*, 2002.

J. L. Hennessy and D. A. Patterson, *Computer architecture: a quantitative approach*, 2011.

Conclusion

- **Demonstrate the key features of Firestorm**
 - Including execution units, latency, throughput, ROB size & register file sizes
- **Accuracy (goal: <10%)**
 - For general cases, SimEng can achieve high accuracy (miss ratio <3%)
 - Worse cases impacted by loop prediction
 - Overall miss-matching ratio 25% of Firestorm

Thank you

Reference

- *Apple unleashes m1*, <https://www.apple.com/uk/newsroom/2020/11/apple-unleashes-m1/>, (2020).
- A. Frumusanu, *Apple announces the apple silicon m1: Ditching x86 - what to expect, based on a14*, <https://www.anandtech.com/show/16226/apple-silicon-m1-a14-deep-dive/2>, (2021).
- *The simulation engine - simeng documentation*, <https://uob-hpc.github.io/SimEng/>, (2021).
- D. A. Patterson and J. L. Hennessy, *Computer organization and design ARM edition: the hardware software interface*, Morgan kaufmann, 2016.
- J. L. Hennessy and D. A. Patterson, *Computer architecture: a quantitative approach*, Elsevier, 2011.
- K. Hwang and N. Jotwani, *Advanced computer architecture: parallelism, scalability, programmability*, vol. 199, McGraw-Hill New York, 1993.
- Cardyak, *Cpu march block diagrams*, <https://drive.google.com/drive/folders/1-Ls0SiqXBCBhMGr87ldA47GETWfjCzQ>, (2022).
- M. Kobayashi, *Dynamic characteristics of loops*, IEEE Transactions on Computers, 33 (1984), pp. 125–132.
- M. De Alba and D. Kaeli, *Path-based hardware loop prediction*, in 4th International Conference on Control, Virtual Instrumentation and Digital Systems, 2002, pp. 29–38.

Backup

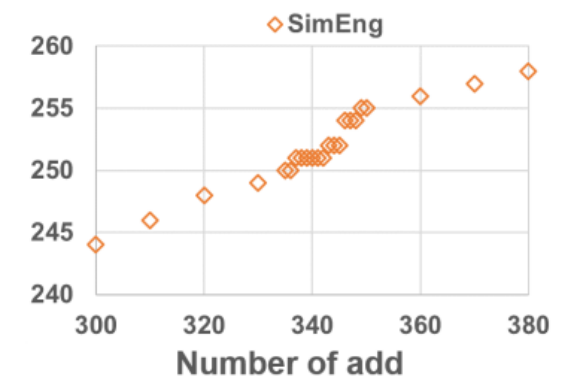
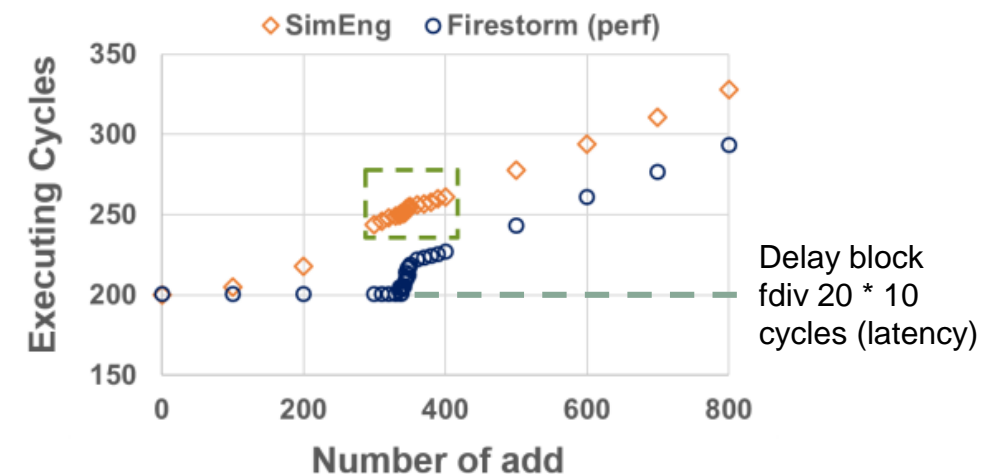
ROB Size & Register File Size

- Buffers/queues in microarchitecture
 - Capable of performing superscalar execution
 - Larger, higher potential performance
- Sizes of buffers
 - Cycles jump as the buffer is filled

	INT register file	FP register file	ROB
My work	344	400	600
Dougall J.	380	432	623
Potential miss-match	9.7%	8.0%	3.7%

- SimEng can match the position of glitch
- However, SimEng gets almost linear curve impacted by fetch rate for loops

For INT register file size,

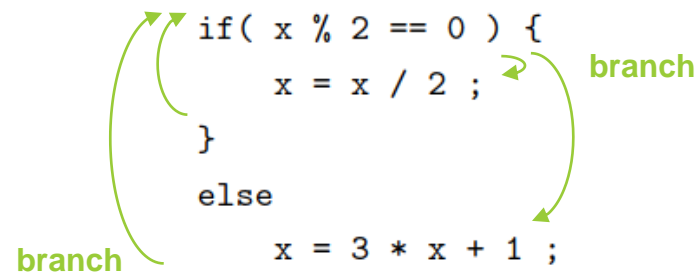


D. Johnson, <https://dougallj.github.io/applecpu/firestorm.html>, (2021).

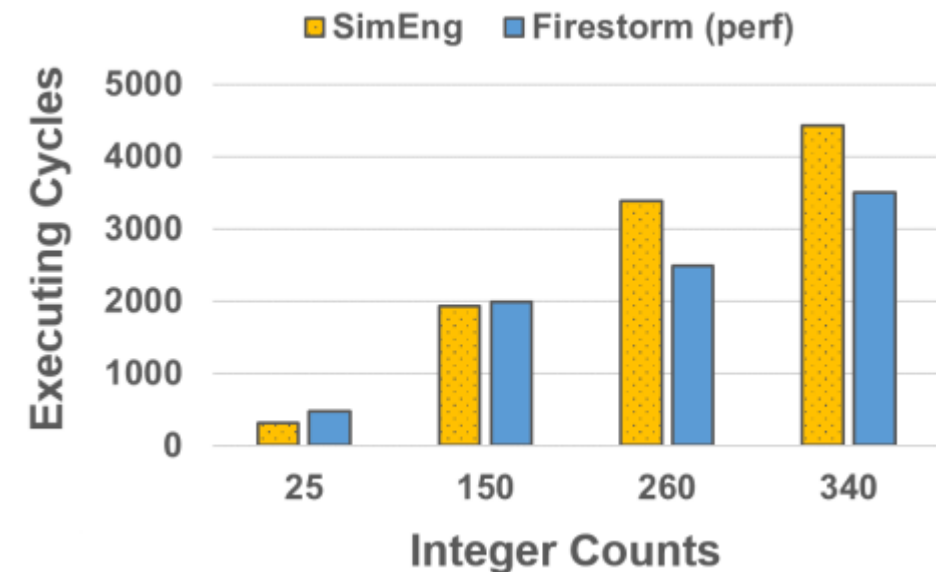
Program: Hailstone Sequence

□ Verifying SimEng with an actual program

- Generating integers w/ rules
 - Every integer takes at least 2 branches



- SimEng match as generating 150 integers
- SimEng's accuracy ↓ as integer counts ↓
 - Overall accuracy: 25%, goal <10%
 - Possibly impacted by loop prediction or branch miss-prediction



Assembly

```
1  .global _start
2  .align 2
3  _start:
4      mov x0, 1
5      mov x1, 2
6      mov x0, #1000
7      mul x0, x0, x0      @ iteration times
8      add x1, x1, #3      @
9      add x1, x1, #3      @ main body
10     add x1, x1, #3      @
11     add x3, x3, #1      @ loop counter
12     cmp x3, x0
13     bne _start+0x10      @ conditional branch
14     mov x0, #0          @ system call
15     mov x8, #94
16     svc #0
```

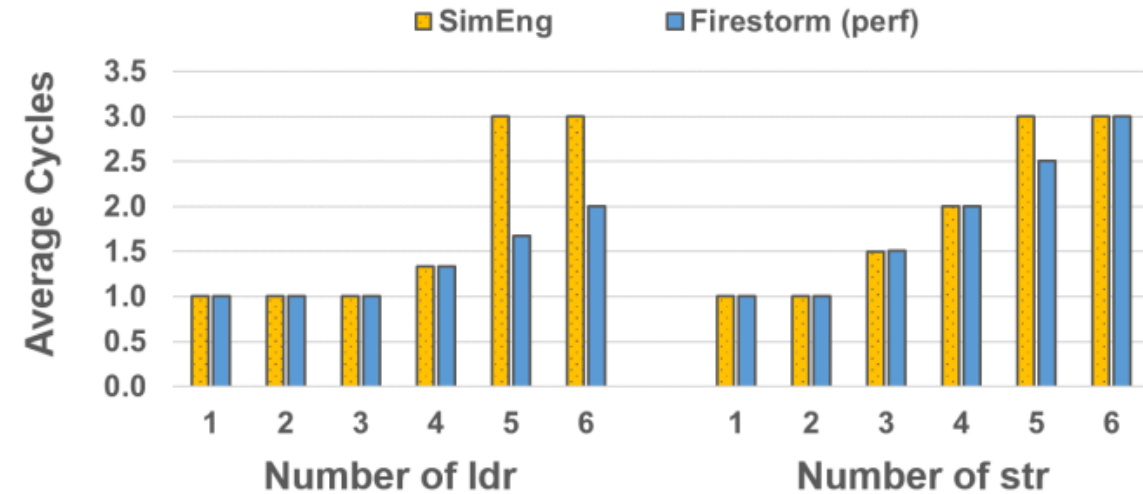
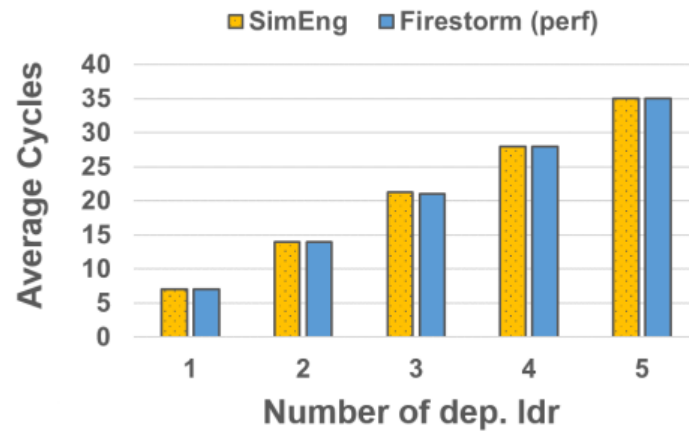
Latency, throughput, and unit counts of opcode groups

Instruction Group List				
Instruction Groups	Opcodes	Latency	Through-put	Unit Counts
INT_SIMPLE_ARITH_NOSHIFT	add	1	1	6
INT_SIMPLE_LOGICAL_NOSHIFT	and	1	1	6
INT_SIMPLE_CMP	cmp	1	1	3
INT_MUL	mul	3	1	2
INT_DIV_OR_SQRT	sdiv	7	2	1
FP_SIMPLE_ARITH_NOSHIFT	fabs	2	1	4
FP_SIMPLE_LOGICAL_NOSHIFT	bic	2	1	4
FP_MUL	fmul	4	1	4
FP_DIV_OR_SQRT	fdiv	10	1	1

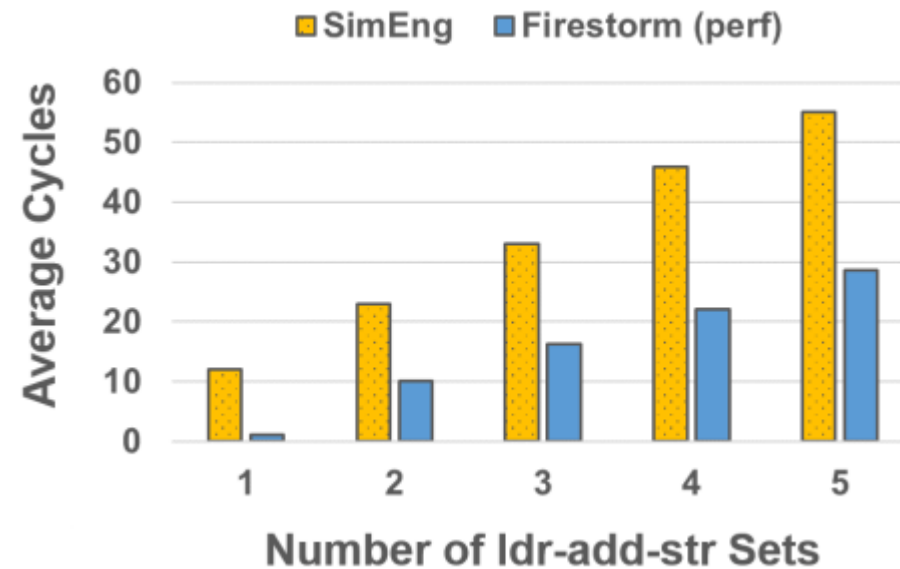
Loop Prediction

```
1  .global _start
2  .align 2
3  _start:
4      mov x0, 1
5      mov x1, 2
6      mov x0, #1000
7      mul x0, x0, x0           @ iteration times
8      add x1, x1, #3           @
9      add x1, x1, #3           @ main body
10     add x1, x1, #3           @
11     add x3, x3, #1           @ loop counter
12     cmp x3, x0
13     bne _start+0x10          @ conditional branch
14     mov x0, #0               @ system call
15     mov x8, #94
16     svc #0
```

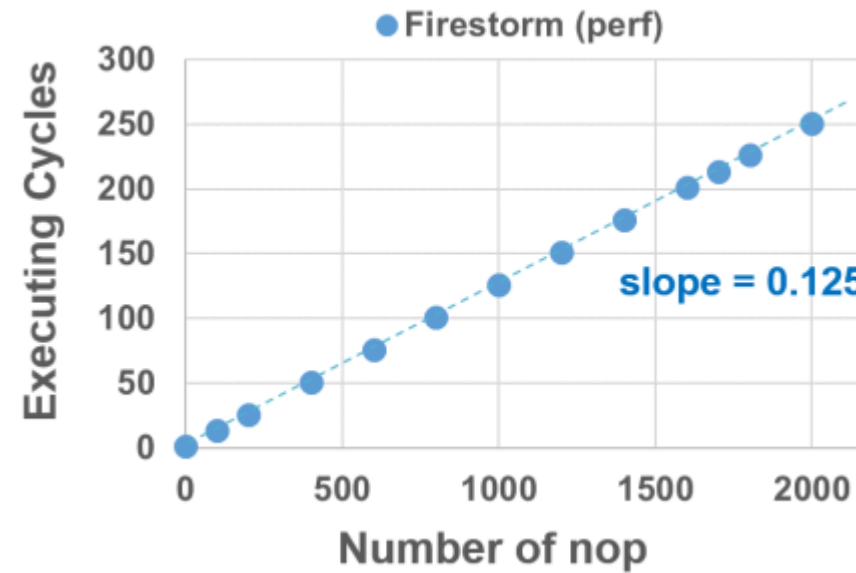
L/S – Latency & Throughput



Ldr-str Depedency

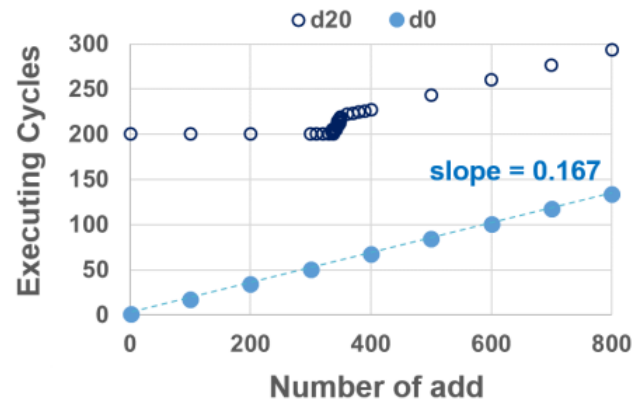


Bandwidth




Perf results of Buffer Size

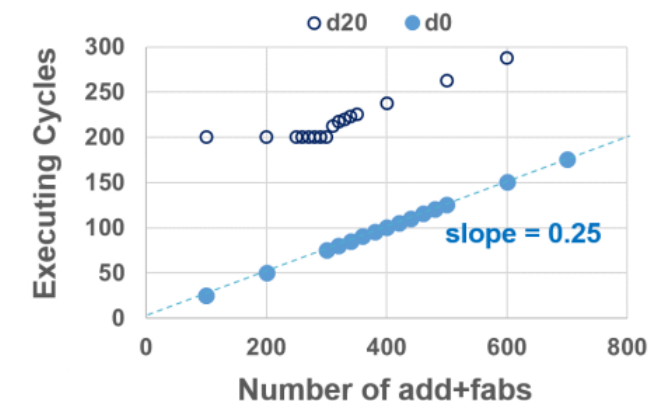
Integer register file



FP register file



ROB



Hailstone Sequence

```
1 .global _start
2 .align 2
3 _start:
4     mov x9, #1
5     mul x9, x9, x9
6     mov x10, #0
7 _loop:
8     mov x7, #3591
9         cmp x7, #1
10        beq _end
11        tst x7, #1
12        beq _even
13        bne _odd
14 _even:
15        mov x6, #2
16        udiv x4, x7, x6
17        mov x7, x4
18        b _loop+0x4
19 _odd:
20        mov x5, #3
21        mul x3, x7, x5
22        add x7, x3, #1
23        b _loop+0x4
24 _end:
25        add x10, x10, #1
26        cmp x10, x9
27        bne _loop
28        mov x0, #0
29        mov x8, #94
30        svc #0
```